

CS 5135/6035 Learning Probabilistic Models

Lecture 1: Course Overview

Gowtham Atluri

August 25, 2018

1 A Gentle Introduction

2 Course Overview

3 Probability and its history

Data Deluge

A Gentle Introduction

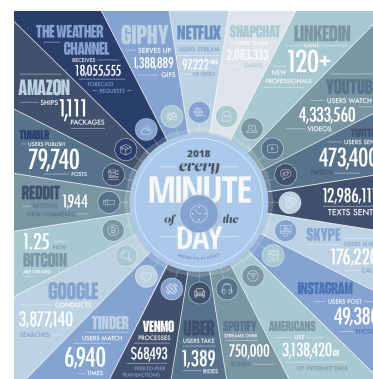


Figure 1: 2018 Internet Data estimates¹


Learning Probabilistic Models

Data Deluge

- Data is being collected at an unprecedented rate
 - More data is being collected every year than ever collected
- Science** - large scale experiments in astronomy, high energy physics, next generation genomics datasets, climate data, etc.
- Business** - e-commerce, online advertising, electronic trading, self-driving cars, etc.
- Society** - Government data, social media, mobile health, public health, crime data, etc.

We seek to harness this data and help discover actionable insights to accelerate science, advance businesses, and address societal problems.

- Mathematical modelling is a process of representing real world problems in mathematical terms in an attempt to find solutions to the problems.
 - Model is typically as a simplification or abstraction of a (complex) real world problem



Learning Probabilistic Models

Types of Models

Deterministic Models


- Output is fully determined by parameter values
- Examples:
 - current through a conductor
 $I = \frac{V}{R}$
 - area of a circular lake:
 $A = \pi r^2$
 - predicting the price of a house based on relevant factors:
 $y = \beta x$

Probabilistic Models


- Poses inherent randomness
- Output can be a prediction that is different even for same parameter values
- Output can be a probability of occurrence of an event
 - E.g., Weather forecasting models: what is the probability it will rain today?

Applications where deterministic models are ineffective

Don't worry about typos




Spelling corrections



Activity detection in smart watches



Speech-to-text conversion



Movie recommendations

Learning Probabilistic Models

"As far as the laws of mathematics refer to reality, they are not certain; and as far as they are certain, they do not refer to reality."

— Albert Einstein

- Source of uncertainty:
 - incomplete/noisy data**
 - not all data can be collected
 - incomplete knowledge**
 - not all functions of a gene are known
 - inherent randomness**
- Probability theory is a mathematical language for **representing and manipulating uncertainty**.



The inevitable reconciliation of **Fortuna** (goddess of chance) and **Sapientia** (wisdom incarnate). 16th century wood engraving.

Learning Probabilistic Models

- Probability theory is a mathematical language for **representing and manipulating uncertainty**.

Advantages of probability models

- They are conceptually simple
 - Probability distributions are used to represent all uncertain unobserved quantities in a model and how they relate to the data.
- Support hierarchical construction
 - Simple probabilistic models of one or a few variables can be used to construct larger, more complex models.
- Easier to understand even complex models
 - The compositionality of probabilistic models makes it much easier to understand the models.

Learning Probabilistic Models

- Given data from a real-world phenomenon
 - How do we choose a suitable family of models to learn?
 - How do we learn a probabilistic model?
 - How do we ensure that the learned model fits the data well?
 - How do we use it for predicting unknown variables?
- Keywords: Learning, Evaluation, and Inference
 - Learning:** Given a set of samples that are known/assumed to be generated from a model, the goal is to determine the parameters of the model.
 - Evaluation:** Given the samples and the learned model, the goal is to determine how well the model fits the data.
 - Inference:** Given a set of model parameters and an observation of some variable(s), the goal is to predict states of other variables.
- Inference is also referred to as *probabilistic reasoning*.

Learning Probabilistic Models

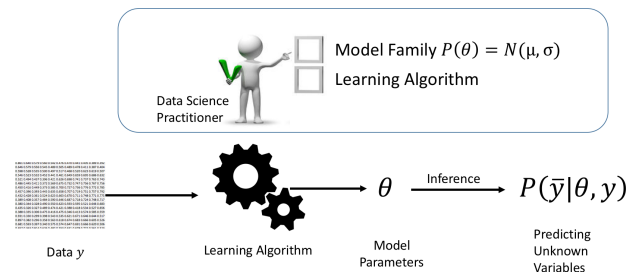


Figure 2: A simplified view

What is in it for Computer Scientists?

- We (in CS) seek to develop systems that can automatically collect necessary data, make decisions and complete tasks.
- Examples:
 - Autonomous vehicles
 - Automated diagnosis of cancers
 - Translating sentences from one language to another
- Computational algorithms are an integral part of probabilistic models that offer effective solutions in these applications.
- We are interested in **learning algorithms** and **inference algorithms**
- To fully appreciate the utility and effectiveness of these algorithms, we need to understand the principles in probability and statistics.

Course Overview

Prerequisites

Courses

This course will build on some foundational concepts from:

- Basic probability and statistics
- Calculus
- Analysis of algorithms
- Discrete mathematics/ Graph Theory

We will do a quick review some of the basic concepts early-on.

Programming Language

- We will use Julia for programming exercises.
 - You are not expected to know Julia before hand.
 - Familiarity with Matlab will be useful.

Course contents

- Descriptive statistics
- Probability foundations
- Discrete probability distributions
- Continuous probability distributions
- Maximum likelihood estimation
- Bayesian approach: Single parameter models
- Bayesian approach: Multi-parameter models
- Connections to non-Bayesian approaches
- Hierarchical models
- Model Checking
- Model Evaluation
- Bayesian computation
- Metropolis-Hastings
- Gibbs Sampling
- Markov chains
- Markov chain Monte Carlo

Learning Outcomes

Upon successful completion of this course, you will be able to...

- determine the suitable probabilistic models for different problem settings
- implement algorithms for learning models based on existing frameworks
- be able to use inference algorithms make inferences in real-world applications
- read previous and existing research literature in this area and critique about the strengths and weaknesses

This course in the larger context

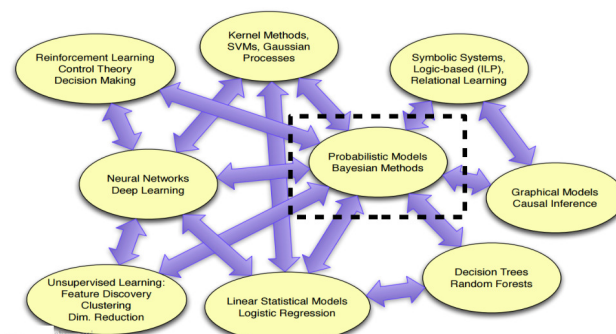



Figure 3: Machine Learning Landscape

Textbook(s)



Grading

Grading scale:

Score breakdown:

- Quizzes: 5%
- Homeworks = 25%
- Midterm exams (2 x 20%): 40%
- Final: 30%

- A: 91-100
- A-: 86-90
- B+: 81-85
- B: 76-80
- B-: 71-75
- C+: 66-70
- C: 61-65

Teaching Methodology


- This class will be taught in a flipped classroom style
 - The lectures videos will be made available prior to the class.
 - These videos will be typically 50 mins/class.
 - Students are expected to watch the lectures prior to participating in the class.
 - There will be a quiz at the beginning of each class.
 - Students will be working in groups on practice questions during the class.
- Please bring your computer to every class.

Flipped Classroom

A flipped classroom is an instructional strategy and a type of blended learning that reverses the traditional learning environment by delivering instructional content, often online, outside of the classroom.^a

^ahttps://en.wikipedia.org/wiki/Flipped_classroom

Programming

- We will be using 
- What is it?
 - a new programming language for scientific computing
 - developed by a group mostly from MIT
 - fully open source, i.e., free
 - convenient syntax for building math constructs like vectors, matrices, etc.
 - largely similar to Matlab
- Some pointers:
 - Tutorial on Julia Syntax [<https://github.com/madeleineudell/intro-to-julia/blob/master/Julia%20Syntax%20Tutorial.ipynb>]
 - Julia Documentation [<https://docs.julialang.org/en/v0.6.0/>]
- We will be using Julia 0.6.4 [<https://julialang.org/downloads/oldreleases.html>]
 - Even though Julia 1.0 was released, many packages we will need are yet to be updated.

Schedule

Date	Topic	Reading
August 28	Introduction and Overview	Probabilistic ML and AI - Zoubin O.
August 30	Descriptive Statistics	Ch 1 Bayesian Computation with R, JA
September 4	Introduction to Probability	Ch 1&2 Intro to Prob & App - RS,LY
September 6	Standard Distributions: Discrete	Ch 4 Intro to Prob & App - RS,LY
September 11	Standard Distributions: Continuous	Ch 5 Intro to Prob & App - RS,LY
September 13	Maximum Likelihood Estimation	Tutorial on MLE - In Jae Myung
September 18	Bayesian Approach I	Ch 2 Bayesian Data Analysis - AG
September 20	Bayesian Approach II	Ch 3 Bayesian Data Analysis - AG
September 25	Bayesian vs. Non-Bayesian approaches	Ch 4 Bayesian Data Analysis - AG
September 27	Hierarchical Models I	Ch 5 Bayesian Data Analysis - AG
October 2	Midterm 1	
October 4		
October 9		
October 11	Fall Reading Days (No Class)	
October 16		
October 18		
October 23		
October 30		
November 1		
November 6	Midterm 2	
November 8		
November 13		
November 15		
November 20		
November 22	Thanksgiving Holiday (No Class)	
November 27		
November 29		
December 4	Final Review	
December 6	Final	

Some guidelines for success in this course

- Assess early-on if the course content is aligned with your interest.
 - A previous student: "Dr. Atluri taught a very math heavy course, and at times it was hard to follow."
- Active participation in the classroom is strongly encouraged.
- If you have difficulty following the course, speak to the instructor.
- Use discussion boards to ask questions. Participate in discussion boards for exam preparation.
- Follow student code of conduct [https://www.uc.edu/conduct/Code_of_Conduct.html].
 - No Plagiarism!

Probability and its history

What is probability?

Definition:

- the chance that a given event will occur ²
- the ratio of the number of outcomes in an exhaustive set of **equally likely** outcomes that produce a given event to the total number of possible outcomes.



- Scenario: Tossing a coin
- A **trial** or **experiment** is one toss of a coin
- Possible outcomes: {Heads, Tails}
- Favourable outcome/event: Heads
- Probability: $\frac{1}{2}$

²<http://www.m-w.com/dictionary/probability>

History of probability - I

The Pascal - Fermat correspondence of 1654

- Cited by historians as the origin of probability theory.
- **The problem of points** - posed by a gambler, Chevalier De Mere.
- **Game:** Two players of equal skill play a game with an ultimate monetary prize. The first to win a fixed number of rounds wins the prize.
- **Scenario:** How should the stakes be divided if the game is interrupted after several rounds, but before either player has won the required number?
- There is *uncertainty* in how the game will unfold.

History of probability - II

- **Existing solutions of the time:** Split the prize in the ratio of points already scored. *Credit for being ahead in the game.*
- Assume player A needs to win a rounds and player B needs b rounds. The game can go at most $a + b - 1$ rounds.
- Number of all possible scenarios 2^{a+b-1} .
- **Pascal and Fermat's solution:** The fair division of the stake will be the proportion of the scenarios that lead to a win by A versus the proportion that lead to a win by B.
 - Takes into account the 'chance' of winning the game.

History of probability - different contexts - I

Other problems from the mid-17th century that required dealing with *uncertainty*

- Annuities
 - Two parties A and B agree that A pays B a lump sum, while B pays back A in annual installments for n years.
 - In the case of life annuity, B pays A a set sum every year until death.
 - To receive \$100 annually from B, how much should B pay A (assuming 5% interest)? *Depends on B's life expectancy.*

History of probability - different contexts - II

Other problems from the mid-17th century that required dealing with *uncertainty*

- Legal system
 - A judge is a '*trier of fact*' i.e., to determine fact based on presented evidence.
 - How much and what kind of evidence was required to produce a 'degree' of conviction in the mind of the judge?
 - Leibniz proposed to calculate the probability of statements of fact in order to determine whether they were true.
 - **Belief proportioned to evidence.**
 - The statements with the highest probability score would be judged to be true.

History of probability - different contexts - III

Other problems from the mid-17th century that required dealing with *uncertainty*

- Plague in London
 - John Graunt, in 1661, computed a '*life table*' that assigns probability of survival to each age using records of birth and death from Church of England.

Table 1. Graunt's Life Table.		
Age Interval	Prop. Deaths in Interval	Prop. Surviving til start of Interval
0-6	0.36	1.00
7-16	0.24	0.64
17-26	0.15	0.40
27-36	0.09	0.25
37-46	0.06	0.16
47-56	0.04	0.10
57-66	0.03	0.06
67-76	0.02	0.03
77-86	0.01	0.01

History of probability - different contexts - III

Other problems from the mid-17th century that required dealing with *uncertainty*

- Plague in London
 - John Graunt, in 1661, computed a '*life table*' that assigns probability of survival to each age using records of birth and death from Church of England.
 - He made a great discovery concerning the prevalence of the plague. A 50 year old person had about as much chance of dying in the next year as did a 20 year old. *Cause of death is not age related, very likely plague related.*
 - The systematic study of quantitative facts about the state – births, deaths, incidence of diseases, emigration, etc – is referred to as political arithmetic.
 - This 17th century political arithmetic was first referred to as '**State-istics**' in late 18th century that is now popularly referred to as '**Statistics**'.

History of probability - books

